

```
In [ ]: import numpy as np
from tensorflow import keras

# Load and flatten MNIST data
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
x_train = x_train.reshape(-1, 28 * 28).astype('float32') / 255.0
x_test = x_test.reshape(-1, 28 * 28).astype('float32') / 255.0

# One-hot encode labels
C = 10
def one_hot(y, C):
    Y = np.zeros((y.shape[0], C), dtype=np.float32)
    Y[np.arange(y.shape[0]), y] = 1.0
    return Y

Y_train = one_hot(y_train, C)
Y_test = one_hot(y_test, C)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 2s 0us/step

```
In [ ]: # Hyperparameters
H = 128 # hidden layer size
rng = np.random.default_rng(0)

# Xavier initialization
W1 = rng.normal(0, 0.02, size=(28 * 28, H)).astype(np.float32)
b1 = np.zeros((H,), dtype=np.float32)
W2 = rng.normal(0, 0.02, size=(H, C)).astype(np.float32)
b2 = np.zeros((C,), dtype=np.float32)
```

Introduces "Hidden Layer" (H) used to learn more complex features. Now a two-layer neural network:

Layer 1 learns features, Layer 2 uses those features to classify digits

```
In [ ]: # Activation and softmax
def relu(u):
    return np.maximum(u, 0.0)

def softmax(z):
    z = z - z.max(axis=1, keepdims=True) # numerical stability
    expz = np.exp(z)
    return expz / expz.sum(axis=1, keepdims=True)

# Forward pass
def forward(X):
    u1 = X @ W1 + b1
    h1 = relu(u1)
    u2 = h1 @ W2 + b2
    yhat = softmax(u2)
    return yhat, (X, u1, h1, u2)

# Accuracy
def accuracy(Y, P):
    return np.mean(np.argmax(Y, axis=1) == np.argmax(P, axis=1))
```

relu() introduces nonlinearity, allowing more complex decision surfaces.

softmax() converts final scores to probabilities

Forward pass:

1. Linear transformation from input to hidden
2. ReLU activation
3. Linear transformation from hidden to output

Intermediate results are used again during backpropagation

```
In [ ]: def train_epoch(X, Y, lr=0.1, batch=128):
    global W1, b1, W2, b2
    idx = rng.permutation(X.shape[0])

    for i in range(0, len(idx), batch):
        j = idx[i:i+batch]
        Xb, Yb = X[j], Y[j]

        # Forward pass
        yhat, (Xc, u1, h1, u2) = forward(Xb)

        # Compute gradients step by step
```

```

du2 = (yhat - Yb) / Xb.shape[0]           # dL/dU2 derivative of the loss w.r.t. output logits.
gW2 = h1.T @ du2                         # dL/dW2 gradient for output weights.
gb2 = du2.sum(axis=0)                     # dL/db2

dh1 = du2 @ W2.T                         # dL/dh1 gradient flowing backward into the hidden layer.
du1 = dh1 * (u1 > 0)                    # dL/dU1 (ReLU derivative) (1 if >0, else 0).
gW1 = Xb.T @ du1                        # dL/dW1 gradient for first layer weights.
gb1 = du1.sum(axis=0)                     # dL/db1

# Gradient descent update
W2 -= lr * gW2
b2 -= lr * gb2
W1 -= lr * gW1
b1 -= lr * gb1

# Train for several epochs
epochs = 8
for ep in range(epochs):
    train_epoch(x_train, Y_train, lr=0.1, batch=128)
    P = forward(x_train[:10000])[0]
    acc = accuracy(Y_train[:10000], P)
    print(f"Epoch {ep+1:02d} - Training accuracy: {acc:.4f}")

```

Epoch 01 - Training accuracy: 0.9092  
 Epoch 02 - Training accuracy: 0.9270  
 Epoch 03 - Training accuracy: 0.9388  
 Epoch 04 - Training accuracy: 0.9466  
 Epoch 05 - Training accuracy: 0.9544  
 Epoch 06 - Training accuracy: 0.9592  
 Epoch 07 - Training accuracy: 0.9628  
 Epoch 08 - Training accuracy: 0.9674

```

In [ ]: # Evaluate on the test set
P_test = forward(x_test)[0]
test_acc = accuracy(Y_test, P_test)
print(f"Test Accuracy: {test_acc * 100:.2f}%")

```

Test Accuracy: 96.33%

Notes: Backpropogation is implemented entirely by hand, every derivative step is shown.

No tensorflow, no autograd, just matrix operations

```

In [ ]: from google.colab import drive
drive.mount('/content/drive')
!jupyter nbconvert --to html '/content/drive/MyDrive/HNR499/HNR499_Model5'

```