

Structured Outline – HNR 499 Final Project

From Equations to Engines: Translating Mathematical Foundations of AI into Motorcycle Part Recognition

Abstract

- One-paragraph overview of the project's arc: theory → implementation → real-world application.
 - Key concept: *AI as math in action* — showing how linear algebra, calculus, and probability underpin machine learning.
 - Summary of workflow:
 - Built and deconstructed MNIST models (in six abstraction stages).
 - Applied the same mathematical reasoning to a YOLOv8 motorcycle part identifier.
 - Highlight core finding: as abstraction decreases, understanding and control increase.
 - End with metrics: MNIST $\approx 98.8\%$ accuracy; YOLOv8 mAP50 ≈ 0.98 .
-

1. Introduction

Goal: establish context, purpose, and relevance.

- Restate original HNR 401 proposal: *bridging the gap between math theory and AI coding practice*.
 - Explain why understanding *how* models learn (not just *that* they learn) matters.
 - Personal motivation: connecting mathematical theory to a mechanical engineering context (Kawasaki Ninja 250 rebuild).
 - Outline of paper structure: math foundations → controlled experiments (MNIST) → applied implementation (YOLOv8).
 - Transitional note: “By the end, the same calculus that adjusts neural weights in a digit recognizer identifies motorcycle parts.”
-

2. Mathematical Foundations of Machine Learning

Purpose: anchor later work in explicit mathematical reasoning.

2.1 Linear Algebra — Representing Data

- Matrices and vectors as the fundamental data structure.
- MNIST digits = 28×28 matrices; each convolution = matrix multiplication.
- Discuss PCA and dimensionality reduction (from Meeting 1 notes).
- *Insert Figure*: example of 3×3 convolution kernel applied to 5×5 patch (from Model 6).

2.2 Calculus — Learning Through Gradients

- Gradient descent as optimization engine; loss minimization.
- Backpropagation = chain rule applied repeatedly.
- *Insert Example*: derivative path from Model 5 where manual backprop was coded.
- Mention role of learning rate and convergence.

2.3 Probability & Statistics — Judging Performance

- Output probabilities via softmax.
 - Metrics: accuracy, precision, recall, mAP.
 - Statistical validation: how overlapping train/val sets inflate performance (ethical transparency).
 - *Insert Table*: summary of metrics definitions.
-

3. Progressive De-Abstraction: From Black Box to Bare Math

Theme: showing how each layer of abstraction removed reveals the underlying mathematics.

3.1 Model 1 – High-Level CNN (Keras)

- Uses predefined layers; achieves 98 %+ accuracy quickly.
- “Push-button AI”: the baseline demonstration of functionality.
- Talking point: efficiency vs. transparency.

3.2 Model 2 – Manual Training Loop

- Introduced `tf.GradientTape()`; explicit loss and gradient computation.
- Now visible: how weights update every batch.
- Reflect: learning feels mechanical, not magical.

3.3 Model 3 – Raw TensorFlow Ops

- Direct control of convolutions, ReLU, pooling, softmax.
- Introduced Glorot initialization; connect to variance control in math.
- Shows complete visibility of matrix operations.

3.4 Model 4 – Pure NumPy Softmax Regression

- Single-layer, fully manual implementation.
- Demonstrates that even a simple linear model can learn via algebraic operations.
- Accuracy $\approx 91\%$; limited but transparent.

3.5 Model 5 – Two-Layer MLP with Manual Backpropagation

- Every derivative step computed manually.
- Connect directly to calculus section — partial derivatives driving optimization.
- Discuss role of hidden layer and nonlinear transformation (ReLU).

3.6 Model 6 – Manual Convolution Visualization

- Hand-built edge and blur filters.
- Illustrates what CNNs “see.”
- *Insert Figures:* horizontal/vertical edge detection, pooling comparison.

3.7 Reflection Table

- Columns: Model #, Abstraction Level, Math Visibility, Accuracy, Key Takeaway.
 - Emphasize: as abstraction \downarrow , insight \uparrow .
-

4. Application Phase — Motorcycle Part Identifier (YOLOv8)

Purpose: demonstrate transfer from theory to engineering practice.

4.1 Motivation & Dataset

- Problem: identifying small, similar components during a rebuild.
- Collected 486 images across 15 classes (list major ones).
- *Insert Figure:* labeled dataset sample (piston vs. valve spring).
- Discuss manual labeling and folder organization (images/train + labels/train).

4.2 Model Architecture

- Explain YOLO (“You Only Look Once”) approach: simultaneous detection + classification.
- Compare to MNIST CNNs:
 - MNIST = single object classification
 - YOLO = multi-object detection with bounding boxes
- Components: backbone \rightarrow neck \rightarrow head.
- *Insert Diagram:* YOLO pipeline annotated with matrix operations.

4.3 Training Process

- Config: 60 epochs, imgsz = 768, batch = 16, AdamW optimizer.
- *Insert Table*: key hyperparameters.
- Note use of automatic mixed precision (AMP) and cosine LR scheduler.
- Mention Google Colab environment, GPU (Tesla T4).

4.4 Results

- mAP50 \approx 0.983, mAP50-95 \approx 0.746.
- Per-class precision/recall values (summarize top and bottom).
- Discuss why overlapping train/val sets raise reported accuracy.
- *Insert Graph*: accuracy and mAP vs epoch (Training vs Validation).

4.5 Comparison to Hand-Drawn Models

- Similarities: convolutional feature extraction, gradient descent, probabilistic output.
 - Differences: YOLO adds spatial localization (bounding boxes + objectness score).
 - Analogy: MNIST identifies “what,” YOLO identifies “what and where.”
 - Reflect on re-encountering the same math in a new domain.
-

5. Ethical & Practical Reflections

- Data ethics: personal vs. public data; small dataset limitations.
 - Reproducibility and transparency — importance of documenting all training settings.
 - Model bias and overconfidence in small sample sizes.
 - Safety/utility of AI in mechanical contexts (assistive vs. replacement).
 - Tie back to HNR program theme: responsible innovation.
-

6. Future Work

- Separate train/val sets, add validation augmentation (flip, rotation, brightness).
 - Expand dataset beyond 486 images.
 - Experiment with mobile exports (TorchScript / ONNX).
 - Potential integration: AR-guided rebuild assistant.
 - Publish as open-source educational dataset.
-

7. Conclusion

- Recap journey: abstract math \rightarrow hands-on implementation \rightarrow practical application.
- Central argument: mathematical literacy enables deeper AI control and creativity.

- Personal reflection: from “seeing AI as magic” to “understanding it as math.”
 - Forward-looking statement: the same process could apply to other engineering systems.
-

8. References & Appendices

- **References:** (MLA-style list from your HNR 401 proposal.)
- **Appendices:**
 - Code excerpts (MNIST and YOLO).
 - Tables of accuracy/mAP per model.
 - Screenshots of training logs.
 - YOLO data.yaml structure for reproducibility.